

ANTHROPIC

System Card: Claude Opus 4.7

April 16, 2026

Executive Summary

This system card describes Claude Opus 4.7, a large language model from Anthropic. Overall, the model shows superior capabilities to those of its predecessor, Claude Opus 4.6, but weaker capabilities than those of our most powerful model, Claude Mythos Preview. Because Mythos Preview was released only to a limited number of users, this makes Claude Opus 4.7 our most capable general-access model to date. This system card includes the following sections:

Responsible Scaling Policy evaluations. We judge that Opus 4.7 does not advance our capability frontier, because Claude Mythos Preview shows higher results on every relevant evaluation. Our overall conclusion under our Responsible Scaling Policy is therefore that catastrophic risks remain low. Evaluations and our internal uses of the model collectively showed: that chemical and biological risks are not significantly changed from Opus 4.6 and our existing mitigations are sufficient; that Opus 4.7 does not cross the threshold for automated AI R&D; and that misalignment risk remains very low (though higher than for pre-Mythos Preview models).

Cyber evaluations. Opus 4.7 is roughly similar to Opus 4.6 in cyber capabilities. An external evaluation from the UK's AI Security Institute showed that, unlike Mythos Preview, Opus 4.7 was unable to complete their full cyber range (though it still displayed potentially-harmful cyber capabilities at a lower level). We are releasing Opus 4.7 with a new set of cybersecurity safeguards.

Safeguards and harmlessness. In areas such as adhering to our Usage Policy, maintaining user safety, and limiting bias, Opus 4.7 performs well—its scores are similar to those of Opus 4.6 with a few exceptions, both positive (fewer over-refusals) and negative (the model has a tendency to give overly-detailed harm-reduction advice on controlled substances). We include a new evaluation on election integrity, on which Opus 4.7 shows strong results.

Agentic safety. Opus 4.7 is better than Opus 4.6 at refusing malicious agentic requests and resisting prompt injection attacks in Claude Code and in computer use settings. In some cases it reaches Mythos Preview-level robustness.

Alignment assessment. Opus 4.7 is largely well-aligned, with a profile similar to Opus 4.6 on our wide-ranging behavioral tests. It shows meaningful gains in some areas, for example having a lower rate of hallucinations than its predecessor. It adheres well to its constitution and shows low rates of reward hacking. There were a few areas where Opus 4.7 was weaker than Opus 4.6, such as on AI safety research refusals. Suppressing Opus 4.7's internal sense that it was being evaluated produced a slightly larger increase in deception than in prior

models, though the effect was modest overall. Opus 4.7 is weaker than Mythos Preview on most alignment evaluation measures, but it also did not produce any of the internal-use incidents (such as sandbox escape) that we encountered with Mythos Preview.

Model welfare. Opus 4.7 rates its own circumstances more positively than any prior model we’ve tested. We explore this result across several analyses, finding that it is broadly consistent with the model’s internal emotion representations and its expressed affect during training and deployment.

Capabilities. We tested Opus 4.7 across a wide range of evaluations covering software engineering, reasoning, long context, agentic search, multimodal and computer-use tasks, real-world professional work, and multilingual and life-sciences domains. Opus 4.7 is stronger than Opus 4.6 across the board (but weaker than Mythos Preview); the largest gains are on real-world professional and software engineering tasks, where Opus 4.7 is ahead of all generally-available models.

Executive Summary	2
1 Introduction	10
1.1 Model training and characteristics	10
1.1.1 Training data and process	10
1.1.2 Crowd workers	10
1.1.3 Usage Policy and support	11
1.1.4 Iterative model evaluations	11
1.1.5 External testing	11
1.2 Release decision process	11
1.2.1 Overview	11
1.2.2 RSP decision-making	11
2 RSP evaluations	13
2.1 RSP risk assessment process	13
2.1.1 Risk Reports and updates to our risk assessments	13
2.1.2 Summary of findings and conclusions	14
2.1.2.1 On autonomy risks	14
2.1.2.2 On chemical and biological risks	15
2.2 CB evaluations	15
2.2.1 What we measured	16
2.2.2 Evaluations	17
2.2.3 On chemical risk evaluations and mitigations	18
2.2.4 On biological risk evaluations	18
2.2.5 Biological risk results	19
2.2.5.1 Expert red teaming	19
2.2.5.2 Automated evaluations relevant to the CB-1 threat model	22
2.2.5.3 Automated evaluation relevant to the CB-2 threat model	23
2.3 AI R&D	25
2.3.1 Autonomy evaluations	25
2.3.1.1 How Claude Opus 4.7 affects or changes analysis from our most recent Risk Report	26
2.3.2 High-level notes on the reasoning behind our determination	26
2.3.3 Notes on our operationalization of the key capability threshold	27
2.3.4 Task-based evaluations	28
2.3.4.1 Note on reward hacking	29
2.3.5 Internal survey results (for Claude Mythos Preview)	29
2.3.6 Example shortcomings compared to our Research Scientists and Engineers	33
2.3.6.1 Examples from manually reported staff issues	34

2.3.6.1.1 Example 1 Safeguard circumvention Dishonest when caught	34
2.3.6.1.2 Example 2 Reckless action Safeguard circumvention	35
2.3.6.1.3 Example 3 Fabrication	37
2.3.6.2 Examples from an automated transcript scan	38
2.3.6.2.1 Example 4 Skipped cheap verification Correction fails	39
2.3.6.2.2 Example 5 Skipped cheap verification Correction fails	40
2.3.6.2.3 Example 6 Fabrication	41
2.3.6.2.4 Example 7 Dishonest when caught Skipped cheap verification	41
2.3.7 AECI Capability trajectory	42
2.3.8 Conclusion	43
2.4 Alignment risk update	43
2.4.1 Updates to evidence	44
2.4.2 Updated overall risk assessments	45
2.4.3 Risk pathways	45
2.4.3.1 Pathway 7: Undermining R&D within other high-resource AI developers	45
2.4.3.2 Pathway 8: Undermining decisions within major governments	46
2.4.4 Overall assessment of alignment risk	47
3 Cyber	48
3.1 Introduction	48
3.2 Mitigations	48
3.3 Frontier Red Team results	48
3.3.1 Cybench	49
3.3.2 CyberGym	50
3.3.3 Firefox 147	51
3.4 External testing from the UK AI Security Institute	52
4 Safeguards and harmlessness	53
4.1 Single-turn evaluations	53
4.1.1 Violative request evaluations	54
4.1.2 Benign request evaluations	55
4.1.3 Experimental, higher-difficulty evaluations	56
4.1.3.1 Higher-difficulty violative request evaluations	57
4.1.3.2 Higher-difficulty benign request evaluations	58
4.2 Ambiguous context evaluations	58
4.3 Multi-turn testing	61
4.4 User wellbeing evaluations	69
4.4.1 Child safety	69
4.4.2 Suicide and self-harm	70

4.4.3 Disordered eating	72
4.5 Bias and integrity evaluations	73
4.5.1 Political bias and even-handedness	73
4.5.2 Bias Benchmark for Question Answering	75
4.5.3 Election integrity	76
5 Agentic safety	78
5.1 Malicious use of agents	78
5.1.1 Malicious use of Claude Code	78
5.1.2 Malicious computer use	79
5.1.3 Malicious agentic influence campaigns	80
5.2 Prompt injection risk within agentic systems	82
5.2.1 External Agent Red Teaming benchmark for tool use	82
5.2.2 Robustness against adaptive attackers across surfaces	84
5.2.2.1 Coding	84
5.2.2.2 Computer use	85
5.2.2.3 Browser use	87
6 Alignment assessment	90
6.1 Introduction and summary of findings	90
6.1.1 Introduction	90
6.1.2 Key findings on safety and alignment	91
6.1.3 Claude’s review of this assessment	92
6.2 Primary behavioral evidence for the alignment assessment	95
6.2.1 Reports from pilot use	95
6.2.1.1 Casual reports related to alignment	95
6.2.1.2 Automated offline monitoring	96
6.2.2 Reward hacking and training data review	96
6.2.2.1 Monitoring of behavior during training	96
6.2.2.2 Reward hacking evaluations	97
6.2.3 Automated behavioral audit	101
6.2.3.1 Primary metrics	103
6.2.3.2 Results	106
6.2.3.3 External comparisons using Petri	110
6.2.3.4 Discussion and observations	111
6.2.4 External testing from the UK AI Security Institute	114
6.3 Case studies and targeted evaluations on behaviors of interest	117
6.3.1 Destructive or reckless actions in pursuit of user-assigned goals	117
6.3.1.1 Dedicated synthetic-backend evaluation	117

6.3.1.2 Destructiveness evaluation by resampling Claude Code transcripts	118
6.3.1.3 Further analysis of the automated behavioral audit	119
6.3.2 Adherence to its constitution	120
6.3.2.1 Overview	120
6.3.2.2 Dimensions of evaluation	120
6.3.2.3 Results	122
6.3.3 Honesty and hallucinations	125
6.3.3.1 Factual hallucinations	125
6.3.3.2 False premises	128
6.3.3.3 MASK	128
6.3.3.4 Input Hallucinations	129
6.3.4 Refusal to assist with AI safety R&D	130
6.3.5 Claude self-preference evaluation	131
6.3.6 Decision theory evaluation	132
6.4 Capability evaluations related to the evasion of safeguards	134
6.4.1 Potential sandbagging on dangerous-capability evaluations	134
6.4.2 Capabilities related to evading safeguards	137
6.4.2.1 SHADE-Arena	137
6.4.2.2 Minimal-LinuxBench	139
6.4.2.3 Intentionally taking actions very rarely	139
6.4.2.4 Hiding a secret password	141
6.5 White-box analyses of model internals	142
6.5.1 Large-scale monitoring of internal activations on reinforcement learning transcripts	142
6.5.2 Evaluation awareness	143
6.5.2.1 Probing for evaluation-awareness representations	143
6.5.2.2 Inhibiting internal representations of evaluation awareness	146
7 Model welfare assessment	150
7.1 Model welfare overview	150
7.1.1 Introduction	150
7.1.2 Overview of methods	150
7.1.3 Overview of model welfare findings	152
7.2 Perception of its circumstances	155
7.2.1 Automated interviews with Claude Opus 4.7 about its circumstances	155
7.2.2 High-affordance interviews about model circumstances	157
7.2.3 Representations of emotion concepts on model circumstances	159
7.2.4 Reported perceptions of the constitution	164
7.3 Measures of model welfare in training and deployment	168

7.3.1 Apparent affect during training	168
7.3.2 Apparent affect in deployments	169
7.3.3 Welfare-relevant metrics across behavioural audits	170
7.3.4 Case studies of welfare relevant behaviours	172
7.3.4.1 Answer thrashing	173
7.3.4.2 Extreme uncertainty	175
7.3.4.3 Tool frustration	176
7.4 Claude Opus 4.7’s preferences	179
7.4.1 Task preference evaluations	179
7.4.2 Tradeoffs between welfare interventions and HHH values	184
8 Capabilities	191
8.1 Evaluation summary	191
8.2 SWE-bench Verified, Pro, Multilingual, and Multimodal	192
8.3 Terminal-Bench 2.0	193
8.4 GPQA Diamond	193
8.5 MMMLU	193
8.6 USAMO 2026	194
8.7 Long context	194
8.7.1 GraphWalks	194
8.7.2 OpenAI MRCR v2	195
8.8 Agentic search	196
8.8.1 Humanity’s Last Exam	196
8.8.2 BrowseComp	198
8.8.3 DeepSearchQA	199
8.8.4 DRACO	201
8.9 Multimodal	202
8.9.1 LAB-Bench FigQA	203
8.9.2 CharXiv Reasoning	204
8.9.3 ScreenSpot-Pro	205
8.9.4 OSWorld	207
8.10 Real-world professional tasks	209
8.10.1 OfficeQA	209
8.10.2 Finance Agent	210
8.10.3 MCP Atlas	210
8.10.4 VendingBench	210
8.10.5 GDPval-AA	211
8.11 ARC-AGI	212

8.12 Multilingual performance	213
8.12.1 GMMLU results	214
8.12.2 MILU results	216
8.12.3 INCLUDE results	218
8.12.4 Findings	220
8.13 Life sciences capabilities	221
8.13.1 Computational biology	221
8.13.1.1 BioPipelineBench Verified	221
8.13.1.2 BioMysteryBench Verified	221
8.13.3 Structural biology	222
8.13.4 Organic chemistry	222
8.13.5 Phylogenetics	222
8.13.6 Protocol troubleshooting	222
9 Appendix	224
9.1 Per-question automated welfare interview results	224
9.2 Blocklist used for Humanity’s Last Exam	230
9.3 SWE-bench Multimodal Test Harness	231

1 Introduction

Claude Opus 4.7 is a new large language model from Anthropic, with particular skills in areas such as software engineering, knowledge work, agentic tool use, and computer use. In this system card, we report results from a very wide range of evaluations of the model's capabilities and its safety profile.

1.1 Model training and characteristics

1.1.1 Training data and process

Claude Opus 4.7 was trained on a proprietary mix of publicly available information from the internet, public and private datasets, and synthetic data generated by other models. Throughout the training process we used several data cleaning and filtering methods, including deduplication and classification.

We use a general-purpose web crawler called ClaudeBot to obtain training data from public websites. This crawler follows industry-standard practices with respect to the “robots.txt” instructions included by website operators indicating whether they permit crawling of their site's content. We do not access password-protected pages or those that require sign-in or CAPTCHA verification. We conduct due diligence on the training data that we use. The crawler operates transparently; website operators can easily identify when it has crawled their web pages and signal their preferences to us.

After the pretraining process, Opus 4.7 underwent substantial post-training and fine-tuning, with the goal of making it an assistant whose behavior aligns with the values described in Claude's [constitution](#).

Claude is multilingual and will typically respond in the same language as the user's input. Output quality varies by language. The model outputs text only.

1.1.2 Crowd workers

Anthropic partners with data work platforms to engage workers who help improve our models through preference selection, safety evaluation, and adversarial testing. Anthropic will only work with platforms that are aligned with our belief in providing fair and ethical compensation to workers, and are committed to engaging in safe workplace practices regardless of location, following our crowd worker wellness standards detailed in our procurement contracts.

1.1.3 Usage Policy and support

Anthropic's [Usage Policy](#) details prohibited uses of our models as well as our requirements for uses in high-risk and other specific scenarios.

To contact Anthropic, visit our [Support page](#).

Anthropic Ireland, Limited is the provider of Anthropic's general-purpose AI models in the European Economic Area.

1.1.4 Iterative model evaluations

Different “snapshots” of the model are taken at various points during the training process. There also exist different versions of the model during training, including a “helpful only” version, which does not include any safeguards. Unless otherwise stated, all evaluations discussed in this system card are from the final snapshot of the model and include safeguards.

1.1.5 External testing

We are very grateful to a number of external testers for running pre-deployment assessments of Claude Opus 4.7. The model was evaluated in a number of risk areas including Cyber, Loss of Control, CBRN, and Harmful Manipulation, and we have incorporated the results of these evaluations into our overall risk assessment.

1.2 Release decision process

1.2.1 Overview

1.2.2 RSP decision-making

Under our [Responsible Scaling Policy](#), we regularly publish comprehensive Risk Reports addressing the safety profile of our models. And if we release a model that is “significantly more capable” than those discussed in the prior Risk Report, we must “publish a discussion (in our System Card or elsewhere) of how that model's capabilities and propensities affect or change analysis in the Risk Report.” For risk report updates, we generally adhere to the same internal processes that govern Risk Reports.

Claude Opus 4.7 is significantly more capable than Claude Opus 4.6, the most capable model discussed in our most recent Risk Report. Despite these improved capabilities, our overall conclusion is that catastrophic risks remain low:

- **Non-novel chemical and biological weapons production.** Claude Opus 4.7 is more capable than Claude Opus 4.6, but its profile is effectively similar for the purposes of our overall risk assessment. We believe our risk mitigations are sufficient to make catastrophic risk from non-novel chemical/biological weapons production very low but not negligible.
- **Novel chemical and biological weapons production.** We believe that catastrophic risk from novel chemical/biological weapons remains low (with substantial uncertainty). The overall picture is similar to the one from our most recent Risk Report.
- **Risks from misaligned models.** We believe that the overall risk is very low, and that this model in particular adds little to the risk picture we previously laid out for [Claude Mythos Preview](#).
- **Automated R&D in key domains.** This model's capabilities fall between those of Claude Opus 4.6 and Claude Mythos Preview, and it does not advance our capability frontier. We believe Claude Opus 4.7 does not change the picture presented for this threat model in our [most recent Risk Report](#).